
PROFILE SUMMARY

- **Site Reliability Engineer** with 7 years of experience keeping high-availability production systems online across payments, edge networking, and SaaS infrastructure, specializing in SLO design, incident response, and chaos engineering.
- Solid technical background across languages (Go, Python), observability tools (Prometheus, Grafana, OpenTelemetry), container orchestration (Kubernetes), infrastructure as code (Terraform), chaos engineering (Gremlin, Chaos Mesh), and cloud (AWS, GCP) with strong fundamentals in Bash, Linux, and TCP/IP fundamentals.
- Deep expertise in SLO-driven reliability engineering, error-budget policies, graceful degradation, and progressive delivery, leveraging methodologies such as blameless postmortems and game days to drive reliable, observable, and recoverable production systems.
- Engaged collaborator working cross-functionally with Engineering, Product, and Support teams in Agile environments, contributing to architecture reviews, error-budget meetings, and post-incident retrospectives with a pragmatic, ownership-first mindset.
- Emerging leader who shares technical excellence and fosters a culture of reliability-first thinking and operational discipline through PR reviews and runbooks, while leading reliability guild sessions and authoring widely adopted production-readiness checklists.

TECHNICAL SKILLS

Observability & Monitoring:	Prometheus, Grafana, OpenTelemetry, Datadog, ELK, PagerDuty
Languages & Scripting:	Go, Python, Bash, SQL
Container Orchestration:	Kubernetes, Docker, Helm, Istio, Argo Rollouts
IaC & Configuration:	Terraform, Ansible, Helm, ArgoCD, Crossplane
Chaos & Performance Testing:	Gremlin, Chaos Mesh, k6, JMeter, Vegeta
Cloud Platforms:	AWS (EKS, RDS, Lambda, Route 53), GCP (GKE, Pub/Sub)
CI/CD & Release:	GitHub Actions, Spinnaker, Argo Rollouts, Flagger
Incident & On-Call:	PagerDuty, Statuspage, Slack workflows, runbook automation

EDUCATION

University of California, Berkeley [B.S. in Computer Science](#)

Berkeley, CA • Sep 2015 - May 2019

WORK EXPERIENCE

Stripe [Senior Site Reliability Engineer](#)

San Francisco, CA • Sep 2021 - Present

- Own end-to-end reliability for payment processing services supporting \$1T+ annual GMV, leading architecture reviews, production readiness, and on-call rotations across 40+ microservices in a polyglot AWS environment.
- Defined and rolled out the SLI/SLO framework for 35 customer-facing services covering availability, latency, and freshness, introducing multi-window burn-rate alerts and error-budget review meetings that cut paging volume by 48% and held all tier-1 services at 99.95% monthly availability.
- Served as incident commander across 18 SEV1/SEV2 outages, coordinating mitigation across Engineering, Support, and Customer Success teams with runbook automation and decision-tree triage, cutting mean time to mitigate from 42 minutes to 11 minutes.
- Built the unified observability platform on Prometheus, Grafana, and OpenTelemetry, defining SLO dashboards, trace-driven alerting, and routing policies across 40+ services, reducing alert fatigue (median pages per week dropped from 34 to 8).

- Reduced team toil from **43%** to **18%** through **certificate-rotation operators**, self-healing DB failover drills, and capacity-rebalancing automation, reclaiming **600+** engineer-hours per quarter of repetitive operational work.
- Owned capacity planning for the **payments tier**, running **k6 load tests**, defining EKS autoscaling policies, and authoring headroom-budget reviews that absorbed a **5x Black-Friday traffic surge** with zero degradation.
- Established the chaos-engineering practice on **Gremlin** and **Chaos Mesh**, running **22** game days covering **region-failure simulations**, dependency outages, and partial Kubernetes node failures, surfacing **47** reliability gaps and validating quarterly DR plans.

Cloudflare **Site Reliability Engineer**

Austin, TX • Aug 2019 - Aug 2021

- Facilitated the blameless postmortem program across **30+** production incidents, driving action-item tracking and a **weekly incident review board**, lifting close rate from **42%** to **88%** within two quarters.
- Defined the production readiness review process for **24** internal services, codifying **canary rollouts**, automated rollback triggers, and safe-deploy gates, reducing change-failure rate from **6.4%** to **1.1%**.
- Owned production operations for **CDN edge caching** including **runbooks**, DR drills, and change management across **190+** POPs globally, partnering with **Security and Networking** to harden against operational risk.
- Worked closely with **Engineering, Product, and Support** teams across **5** product surfaces to negotiate **error-budget policies**, paging severity thresholds, and incident-response standards, authoring **9** reliability RFCs that shaped the org's reliability-first roadmap and onboarding **12** new SREs.