

PROFILE SUMMARY

- Data Engineer with 6 years of experience designing and operating high-throughput data platforms across SaaS analytics, fintech transaction systems, and product-event pipelines, specializing in dimensional modeling, low-latency streaming, and end-to-end data quality.
- Solid technical background across batch processing (Spark, dbt), streaming (Kafka, Flink), warehouses (Snowflake, BigQuery), orchestration (Airflow, Dagster), and cloud ecosystems (AWS, GCP) with strong scripting fundamentals in Python and SQL.
- Deep expertise in dimensional modeling, lakehouse architecture, data contracts, and CDC ingestion patterns, leveraging methodologies such as Kimball star schemas and medallion architecture to drive trustworthy, queryable, and auditable datasets.
- Engaged collaborator working cross-functionally with Analytics, ML, and Product teams in Agile environments, contributing to data-modeling reviews, SLA negotiations, and stakeholder workshops with a pragmatic, outcome-oriented mindset.
- Emerging leader who shares technical excellence and fosters a culture of data quality and cost discipline through code reviews and runbooks, while leading data-platform working groups and authoring widely adopted contract templates.

TECHNICAL SKILLS

Languages & Scripting:	Python, SQL, Bash, Scala (basic)
Warehouses & Lakehouses:	Snowflake, BigQuery, Redshift, Databricks
Processing & Transformation:	Spark, dbt, Flink, Kafka, Kinesis
Orchestration:	Airflow, Dagster, Prefect, dbt Cloud
Storage Formats & Lakes:	S3, GCS, Iceberg, Delta Lake, Parquet, Avro
Data Quality & Lineage:	dbt tests, Great Expectations, Soda, OpenLineage
Cloud Platforms:	AWS (S3, Glue, EMR, Lambda, IAM), GCP (BigQuery, Dataflow, Pub/Sub)
DevOps & Tooling:	Terraform, GitHub Actions, Docker, Kubernetes, Datadog, Git

EDUCATION

University of California, Berkeley B.S. in Computer Science Berkeley, CA • Sep 2015 - May 2019

WORK EXPERIENCE

Stripe Data Engineer Austin, TX • Aug 2022 - Present

- Own the analytics data platform supporting hundreds of internal dashboards and ML training pipelines, leading end-to-end design and operation across pipeline reliability, data modeling, and cost performance within a modern cloud-native data stack.
- Built and maintained a fleet of 120+ ELT pipelines using dbt and Airflow, moving transactional and event data from Kafka, Postgres, and internal APIs into Snowflake, with parameterized DAG templates that cut new-pipeline onboarding time from 3 days to 4 hours.
- Designed a star-schema warehouse in Snowflake using dbt with SCD Type 2 dimensions, accumulating-snapshot fact tables, and contract-tested staging layers, enabling self-serve analytics for 200+ internal users while keeping query latency under 3 seconds on critical dashboards.
- Optimized Snowflake storage and compute costs through clustering keys, micro-partition pruning, and warehouse auto-

suspend policies, reducing monthly compute spend by **38%** while improving **p95** query latency by **44%** across reporting workloads.

- Migrated **80+** scheduled jobs from **cron + Lambda** to **Airflow** with the **TaskFlow API**, **custom sensors** for data-availability alerts, and DAG-level retry policies, eliminating silent failures and improving on-time delivery rate from **82%** to **99.4%**.
- Stood up a real-time event ingestion pipeline using **Kafka**, **Flink**, and **Iceberg** with exactly-once semantics, watermark-based windowing, and stateful aggregations, delivering fresh-within-**60 seconds** metrics to product teams across **12+** event topics.
- Implemented data quality at every layer using **dbt tests**, Great Expectations suites, and Soda anomaly detection on **180+** critical tables, raising test coverage from **22%** to **91%** and catching **15-20** data-quality regressions per quarter before they reached production dashboards.

Twilio **Data Engineer**

San Francisco, CA • Jul 2019 - Aug 2022

- Ingested data from **60+** source systems including APIs, OLTP databases, S3 dumps, and SaaS connectors using **Fivetran**, custom Python connectors, and CDC via **Debezium**, unifying transaction, customer, and event data into a single normalized warehouse layer with **99.7%** freshness SLA.
- Managed **Postgres**, Redshift, and S3-based data lake tiers across **development, staging, and production** environments, implementing **partitioning**, vacuum scheduling, and storage tier transitions that reduced average query cost per TB by **42%**.
- Provisioned **AWS** data infrastructure using **Terraform** modules, set up CI/CD for **dbt** and **Airflow** code via **GitHub Actions**, and containerized Python ETL jobs with Docker, cutting infrastructure provisioning time from days to **under 30 minutes** and reducing deployment failures by **55%**.
- Implemented **column-level access** controls, PII tagging, and data lineage via **OpenLineage**, with **GDPR**-compliant deletion workflows and pipeline SLA dashboards in **Datadog**, surfacing **99.5%** uptime against published SLAs across financial reporting datasets.